

# VEDANT PUNGLIYA

+1 (515) 815-0442 — [vedantpungliya29@gmail.com](mailto:vedantpungliya29@gmail.com) — [LinkedIn](#) — [GitHub](#) — [Website](#)

## EDUCATION

Iowa State University

Ames, IA

Bachelor of Science, Computer Science

Expected May 2026

GPA: 3.99/4.00 — Minors: Cyber Security Engineering, Data Science, Economics

Coursework: Data Structures & Algorithms, Deep Learning, Full Stack Development, Database Management Systems

## TECHNICAL SKILLS

Languages: Python, Java, C++, SQL, R, Bash, JavaScript (Node.js), HTML/CSS

AI/ML: PyTorch, TensorFlow, LLMs, RAG, LangChain, Vector Databases (FAISS), Model Context Protocol (MCP)

Cloud/DevOps: AWS (EC2, ECS, Fargate, S3, IAM, VPC), Docker, Kubernetes, Jenkins, CI/CD, Terraform, Git

Backend/Web: Flask, FastAPI, Spring Boot, REST APIs, Microservices, Redis, RabbitMQ, PostgreSQL

## PROFESSIONAL EXPERIENCE

PrecisePK

San Diego, CA

Software & Algorithm Developer

Sep 2025 – Dec 2025

- Implemented FOCE-I (First Order Conditional Estimation) algorithms in **Python** and **C++** for nonlinear mixed-effects pharmacokinetic modeling, improving convergence stability by 15%.
- Architected an **AWS**-based HPC pipeline using ECS, CodeBuild, S3, and DynamoDB for scalable job execution.
- Optimized 1/2/3 compartment models to enhance parameter estimation accuracy for clinical studies.

2GoProducts

San Diego, CA

Generative AI & Software Developer

Jun 2025 – Dec 2025

- Engineered multi-agent microservices using Model Context Protocol (MCP) for fault-tolerant orchestration.
- Built a Retrieval-Augmented Generation (**RAG**) text-to-SQL system with FAISS and SQLGlot, enabling accurate natural language database analytics.
- Deployed **AWS** infrastructure (ECS Fargate, API Gateway, ALB) with CI/CD pipelines supporting 11 microservices.
- Designed a multi-tenant GenAI analytics platform with Slackbot integrations, reducing analyst workload by 60%.

Iowa State University

Ames, IA

Undergraduate AI/ML Research Scientist

Aug 2024 – Jun 2025

- Researched **Code Language Models** interpretability using latent-space analysis and saliency methods to detect hallucinations.
- Developed activation steering guardrails (**Keyword + LLM as Judge**), reducing insecure code generation by 65%.
- Led concept ablation studies to validate model robustness; results published in peer-reviewed workshop proceedings.

WeekendAI

Pune, India

AI Engineering Intern

May 2024 – Aug 2024

- Developed autonomous multi-agent systems using **LangChain** and CrewAI, improving pipeline efficiency by 90%.
- Engineered a customer support chatbot grounded in company data using advanced **RAG**, reducing manual ticket resolution time by 40%.
- Implemented real-time system observability dashboards using **Prometheus** and **Grafana** alongside **Phoenix AI**.

## SELECTED PROJECTS

- BioSignify: ML Authentication Platform** — Built Flask microservices with blockchain and vector databases to process 10K+ authentications/day with 99.2% accuracy and sub-200ms latency.
- Intelligent NLP-to-SQL Engine** — Developed a Transformer-based text-to-SQL system with schema inference, optimizing 500+ daily enterprise queries.
- TravelBuddy** — Engineered a Spring Boot platform with an LLM-based recommendation engine and CI/CD deployment; achieved 99.99% uptime.

## CERTIFICATIONS & HONORS

[AWS Certified Solutions Architect – Associate](#)

Nov 2025

Honors: Phi Beta Kappa, IEEE Eta Kappa Nu (HKN), Phi Kappa Phi, President's Excellence Scholarship

Leadership: CS Student Advisory Council; Workshop Lead – CS & Software Engineering Club